
DashVMC: Real-Time Discrete World Model Control in Geometry Dash

Florent Tardieu
INSA Rouen Normandy
florent.tardieu@insa-rouen.fr
[Project page](#) [Code](#)

Abstract

World-model agents are usually evaluated in simulators that can wait for the policy; live games impose the opposite constraint, requiring capture, prediction, and action before the next frame. We present *DashVMC*, a discrete Vision–Model–Controller system that plays Geometry Dash from screen pixels. An FSQ tokenizer maps 64×64 Sobel frames to 8×8 token grids, a block-causal transformer predicts action-conditioned transitions, and a lightweight actor–critic is initialized by behavioural cloning (BC) and improved with PPO in latent rollouts. The frozen policy more than doubles BC mean survival on each of the first two official levels in controlled live evaluation, with a smaller gain on the harder third level. The optimized capture-to-action loop averages 16.3 ms over 5,000 frames on an RTX 2060, below the 33.3 ms budget of the 30-FPS data cadence. Beyond control, the same recurrent dynamics support interactive, open-ended visual level continuations that often remain plausible for extended stretches before autoregressive errors accumulate. Together, these results show that a compact discrete world-model controller can support imagined policy improvement, generative rollouts, and live reactive control within a consumer-GPU frame budget.

1 Introduction

World models learn the consequences of actions and can move policy improvement from costly environment interaction into imagination [1]. Discrete approaches such as IRIS encode images into tokens and model their evolution with an autoregressive transformer [2]; DIAMOND instead shows that preserving fine visual detail with a diffusion model can improve control [3]. Most evaluations, however, run inside an emulator or a learned simulator whose clock effectively waits for inference. An agent controlling a live, non-paused application must also satisfy a wall-clock contract: screen capture, preprocessing, state estimation, policy inference, and action dispatch must finish before the action becomes stale.

Geometry Dash is a small but unforgiving test of this contract. The avatar moves continuously, the action space is binary, and a mistimed jump causes immediate death. Official levels are deterministic, but the agent observes only captured pixels and acts through keyboard events; it does not receive emulator state during control. This combination isolates two practical questions: can a compact visual world model support policy learning, and can the resulting perception-to-action path run within a consumer-GPU frame budget?

DashVMC follows the Vision–Model–Controller decomposition [1]. The vision module is a reconstruction-anchored Finite Scalar Quantization (FSQ) tokenizer [4] over Sobel edge maps. The model is a transformer that interleaves actions with frame blocks and predicts the next token grid in parallel. The controller combines the current token grid with the transformer’s temporal state, receives a short BC warm start, and is then optimized by PPO [5] entirely in frozen-model rollouts.

At deployment, the decoder is absent and the system executes one reactive pass per frame; expensive policy optimization has already happened offline.

This is a systems co-design contribution rather than a new generic world-model primitive. The compact tokenizer and parallel next-frame predictor bound inference cost; explicit action conditioning permits counterfactual rollouts; BC avoids an expensive random-policy transient; and GPU-resident state plus captured execution remove avoidable synchronization from the live path. Our contributions are:

1. an end-to-end discrete world-model controller for a non-paused game, with visual dynamics learned from captured trajectories and a BC-initialized policy improved in imagined rollouts seeded from recorded contexts, then deployed from live screen pixels;
2. controlled live comparisons of the frozen PPO policy against BC and no-op controls across the first three official levels; and
3. an end-to-end latency profile on consumer hardware, matched action interventions, and qualitative interactive rollouts demonstrating open-ended visual level continuation together with its recurrent failure modes.

2 Related work

World models and imagination. The original Vision–Model–Controller pipeline compressed observations with a VAE, modeled dynamics recurrently, and optimized a small controller in dreams [1]. IRIS instead learns a discrete image language and composes it over time with a transformer, enabling policy learning entirely in imagination [2]. Δ -IRIS reduces the sequence cost of this template by separating stochastic changes into discrete tokens from continuous context tokens [6]. DIAMOND removes discrete compression from the dynamics model and uses diffusion to preserve control-relevant visual details [3], while DreamerV3 demonstrates broad control with categorical recurrent state-space latents [7]. Contemporaneous MIRA scales action-conditioned latent diffusion to four-player Rocket League and combines factorized space–time RoPE, QK-normalized attention, SwiGLU feed-forward layers, and AdaLN action conditioning for real-time video generation on a B200 [8]. DashVMC shares the factorized positional abstraction but instead uses discrete FSQ tokens and one-pass next-grid prediction to close an end-to-end capture-to-action loop on consumer hardware. Its transformer also uses the action-conditional contrastive auxiliary objective introduced by TWISTER [9]. MineWorld likewise interleaves discrete visual and action tokens and predicts each frame’s spatial tokens in parallel, while Matrix-Game 2.0 uses few-step diffusion for streaming action-conditioned video [10, 11]. Those systems target interactive video generation; DashVMC instead uses the learned dynamics to optimize a policy and then closes the control loop through the original live game. Geometry Dash has also been controlled from screenshots [12], but that earlier framework injected a fabricated timestep to discretize the application, whereas our protocol leaves the game clock running.

Compact tokenization and policy initialization. FSQ replaces learned vector-quantization codebooks with bounded scalar levels, avoiding commitment losses and codebook collapse while exposing an integer coordinate lattice [4]; we use the improved iFSQ bounding function [13]. For policy initialization, demonstration-based RL has repeatedly shown that supervised behaviour can reduce the expensive early interaction phase [14, 15]. Here demonstrations only initialize the controller: the final policy is selected after PPO improvement in the learned environment.

Representation anchors and alternatives. Joint-embedding prediction can avoid spending capacity on high-magnitude input detail that is irrelevant or unpredictable from context [16], and LeWorldModel demonstrates stable end-to-end predictive learning from pixels with continuous latents [17]. Our observation is instead a deterministic Sobel rendering with much nuisance detail already removed, making reconstruction a comparatively inexpensive anchor for a shared discrete token space. The pre-freeze joint-training diagnostic in Section 5.4 is therefore a constrained FSQ-token adaptation, not a faithful JEPA reproduction or evidence against joint embedding. Likewise, Gaussian Quant reports stronger image rate–distortion results than FSQ on large image benchmarks [18]; FSQ is used here for stability, explicit lattice structure, and measured deployment cost rather than claimed tokenizer optimality.

3 DashVMC

The three modules are trained sequentially: V and M are frozen before controller fitting, and the selected controller is frozen for deployment. Figure 1 separates the imagined PPO path from the live reactive path.




Figure 1: DashVMC policy improvement and deployment. **(a)** PPO rollouts start from a recorded $K = 4$ token-action context. The frozen world model predicts $\hat{\mathbf{z}}_{\text{next}}$ and p_{death} while returning context feature $\mathbf{h}$; the predicted grid and controller-selected action are appended to the rolling context. Four imagined burn-in steps replace the recorded seed before transitions are retained, and PPO updates only C . **(b)** During live control, V encodes captured frames and M performs context prefill only to produce $\mathbf{h}_t$. The controller receives the current grid $\mathbf{z}_t$ directly together with $\mathbf{h}_t$ and dispatches the keyboard action; neither next-grid prediction nor decoding is run.

3.1 Visual tokenization

Each RGB capture is cropped, converted to grayscale, filtered with Sobel edges, and resized to 64×64 . A convolutional autoencoder maps the result through three stride-2 stages to a four-channel 8×8 latent map. Each spatial vector is independently quantized with iFSQ levels [8, 5, 5, 5], using the bounding function $2 \text{sigmoid}(1.6z) - 1$, and yields 64 token IDs from 1000 implicit codes. The tokenizer is trained on adjacent frame pairs with pixel-MSE reconstruction plus temporal slowness and latent-uniformity regularizers [19]. After selection, the encoder is frozen and the decoder is retained only for visualizing dreams.

3.2 Action-conditioned dynamics

The dynamics model has 8 transformer layers, width 384, and 8 attention heads. Each frame block contains 64 visual tokens and one ALIVE/DEATH status token; action tokens are inserted between blocks. Given $K = 4$ observed frame-action pairs, a final masked block requests the next frame. The resulting sequence has $K(65 + 1) + 65 = 329$ positions; every target-frame position receives a learned [MASK] embedding, so no target token is exposed during training. Attention is bidirectional within a frame and causal across frame/action blocks, so all next-frame positions are predicted in

one forward pass without seeing their targets. Factored 3D rotary embeddings encode row, column, and time, with spatial and temporal bases 10 and 10,000; status and action tokens use the grid centre coordinates [20]. The status logit supplies the rollout termination probability, and the final context state $\mathbf{h}_t \in \mathbb{R}^{384}$ summarizes recent dynamics for the controller. In the reported interleaved sequence, this state is taken from the most recent action-token position. Because death frames are oversampled during training, the sigmoid output is used operationally as a termination score rather than interpreted as a calibrated environment probability.

We add TWISTER’s action-conditional contrastive predictive-coding loss with weight 0.1 [9]: for each future offset $k \leq K$, an MLP predicts $\mathbf{h}_{t+k}$ from $\mathbf{h}_t$ and the intervening actions and is scored with in-batch InfoNCE. We also independently replace 5% of context tokens uniformly and 5% with adjacent FSQ codes. The token loss uses focal modulation [21] and a fixed local FSQ-coordinate target. For target code i with lattice coordinate $\mathbf{c}_i$,

$$q(j | i) = \begin{cases} 1 - \varepsilon, & j = i, \\ \varepsilon \frac{\exp(-\|\mathbf{c}_i - \mathbf{c}_j\|_2^2 / 2\sigma^2)}{\sum_{\ell \neq i} \exp(-\|\mathbf{c}_i - \mathbf{c}_\ell\|_2^2 / 2\sigma^2)}, & j \neq i, \end{cases} \quad (1)$$

with $\varepsilon = 0.1$ and $\sigma = 0.9$. Unlike uniform label smoothing [22], this assigns the smoothing mass locally on the FSQ lattice. Writing $H(q, p) = -\sum_j q(j | i) \log p_j$, focal modulation is applied once to the aggregate soft-target loss as $(1 - e^{-H})^2 H$, rather than independently per class; the output projection is tied to the token embedding. Decoder perturbations motivated this target: adjacent FSQ coordinates reconstructed nearly identical patches, with visual error increasing with coordinate offset. In an archived same-architecture comparison on an earlier 256-width model, replacing uniform smoothing with SLS improved validation token accuracy (32.85% to 33.66%) and CPC loss (0.345 to 0.238) at an essentially unchanged train–validation gap. This single-run development comparison supports the local choice; a later matched IRIS/Pong diagnostic does not support a general claim (Section 6).

3.3 Controller and latent policy improvement

The $\sim 45\text{K}$ -parameter actor–critic embeds the current 8×8 token grid, processes it with two convolutions and max pooling, concatenates the resulting 256-dimensional feature with $\mathbf{h}_t$, and uses zero-initialized heads to predict jump probability, value, and the next eight actions. The controller is first fit to expert actions by BC, then optimized with PPO in autoregressive rollouts from the frozen world model.

Each PPO iteration starts from recorded four-frame contexts and generates 512 imagined episodes of at most 45 steps. The first four imagined transitions refresh the rolling context and are excluded from PPO updates; retained transitions therefore use an entirely model-generated four-frame context. The 45-step horizon roughly covers obstacles visible in the recorded seed. Beyond it, they have usually left the crop, so the model often enters a physically valid but uninformative empty segment; later obstacles must also be generated rather than seed-observed, progressively shifting controller inputs away from encoded real frames. We stop at this boundary to learn consequences of observed obstacles without relying on invented challenges. Rollouts terminate when predicted death probability exceeds 0.5. Reward is +1 per survived step and -0.25 per jump, preventing an always-jump local optimum. We use GAE with $\gamma = 0.995$, $\lambda = 0.95$, PPO clip 0.2, and auxiliary-action loss weight 0.1.

3.4 Reactive deployment path

Deployment is fixed at the 30-FPS data cadence. Captured images enter pinned host memory and are transferred asynchronously; token context remains in a GPU-resident ring buffer. At each step, the transformer runs context prefill only to obtain $\mathbf{h}_t$; its masked next-grid prediction phase is not executed. The FSQ encoder uses `torch.compile`, and transformer context encoding is captured in a CUDA Graph after warm-up. The decoder is omitted, and a keyboard update is issued only when the binary action changes. The loop sleeps after action dispatch to maintain cadence, but that deliberate sleep is excluded from capture-to-action latency. This is deliberately reactive rather than an inference-time planning system: online search or trajectory optimization would have to share the same 33.3-ms budget with capture and dispatch.

4 Experimental setup

Data and training. The corpus contains 4,264 base gameplay episodes (about 250K frames), including deliberate failures and 36 expert trajectories or segments. The base episodes span primarily official Levels 1–7 and multiple community-created levels. Episodes containing mechanics outside the selected model’s scope, including gravity inversion, were excluded to keep the learned dynamics tractable at this scale. The expert subset comprises 33,153 frames (18.95 minutes). Capture metadata do not identify every segment’s in-level start state or completion status, so expert segments are not treated as verified full-level clears. Episodes are split once, 90/10, by base identifier and source with seed 42; vertical shifts of ± 2 and ± 4 pixels inherit the base assignment and expand the tokenized corpus to 21,320 episodes. The frozen encoder retokenizes the corpus before dynamics training. All stages use seed 42; primary training is sequential in BF16 on one A100, while live inference is measured on an RTX 2060. Table 1 records the optimization and selection details that materially determine the frozen system.

Table 1: Training and checkpoint selection. Tokenizer/dynamics learning rates use cosine decay after any stated warm-up.

Stage	Optimization	Selection
Tokenizer	Adam; 1,000 epochs; batch 2,048; cosine LR from 10^{-3} to 10^{-5} ; slowness 0.1; uniformity 0.01; validation every 10 epochs.	Min. loss (epoch 920).
Dynamics	AdamW; 200 epochs $\times$ 500 steps; batch 512; LR $2 \times 10^{-3} \rightarrow 5 \times 10^{-5}$ after 5-epoch warm-up; weight decay 0.01; dropout 0.1; max grad. norm 1.0; death frames oversampled $5\times$.	Max. death F1 (epoch 139).
BC	AdamW; 50 epochs; batch 512; LR 10^{-3} ; weight decay 10^{-4} ; positive jump weight 1.5.	Min. loss (epoch 9).
PPO	Adam; 15,000 iterations; 512 rollouts; four update epochs; minibatch 512; LR 10^{-4} ; entropy/critic weights 0.01/0.5; ratio clip 0.2; max grad. norm 0.5.	Max. development survival (iteration 3,250).

Checkpoint and live evaluation. PPO checkpoint selection uses survival on 512 fixed latent development contexts every ten iterations. Because training seeds may originate from the same episode collection and selection is adaptive, this is explicitly a development metric, not held-out validation. For live evaluation, all checkpoints are frozen and Auto-Retry is enabled. Auto-Retry only standardizes restarts and supplies no observation to the policy; the policy receives neither a level identifier, elapsed time or frame index, level progress, nor game memory. Live actions are deterministic: jump is selected exactly when the controller output exceeds 0.5. We retain 99 scored PPO attempts, 100 BC attempts, and 10 no-op attempts on each of the first three official levels. The deterministic no-op diagnostic is repeated ten times because it always reaches the first obstacle; the repeats quantify capture and action-timing jitter, with a standard deviation of 0.7 frames on every level. Each invocation’s manually resumed synchronization episode is excluded; game memory is used only to delimit alive/dead episodes, never as policy input. The endpoint is acted frames survived, and reported PPO intervals are 95% percentile-bootstrap confidence intervals for the policy mean using 50,000 resamples. The original selected BC file was not archived, so the live control uses an epoch-9 rerun whose validation accuracy is close, but not bit-identical, to the archived trace (79.85% versus 79.76%).

Latency and continuation protocols. Latency is measured on 5,000 consecutive frames of the optimized deployment path, from capture through action update, with CUDA synchronization only at the end-to-end boundary. For the qualitative intervention, a human-controlled rollout supplies a real four-frame context before a spike. Greedy futures start from identical reconstructed tokens and differ only in the first action; all subsequent actions are idle. The display-only HUD is removed before re-encoding the context.

5 Results

5.1 Frozen-system and live-control results

All component metrics below are reported at their checkpoint-selection points rather than as independent maxima. The selected tokenizer reaches 3.894×10^{-4} per-pixel validation MSE (34.10 dB PSNR), obtained by dividing its per-frame reconstruction SSE of 1.595 by 64^2 . The transformer reaches 29.74% token accuracy and death precision/recall/F1 of 0.7337/0.8654/0.7941; low exact-token accuracy can coexist with useful local predictions in a 1,000-class grid, but it also bounds dream reliability. Selected BC accuracy is 79.76%, and the selected PPO checkpoint survives 29.71 of 45 steps on the fixed latent development contexts.

Table 2: Live acted-frame survival (mean $\pm$ standard deviation). Attempts are independent restarts of one frozen policy; n is per level. BC uses the epoch-9 rerun described in the protocol.

Policy	n	Level 1	Level 2	Level 3
No-op	10	46.2 $\pm$ 0.7	63.4 $\pm$ 0.7	41.5 $\pm$ 0.7
BC	100	130.4 $\pm$ 91.1	119.7 $\pm$ 71.5	45.2 $\pm$ 9.0
PPO	99	279.4 $\pm$ 48.0	262.9 $\pm$ 121.2	63.8 $\pm$ 27.1

Table 2 shows that PPO adds 149.0 and 143.2 mean frames over BC on Levels 1 and 2, more than doubling survival; the PPO mean’s 95% bootstrap intervals are [270.2, 289.0] and [239.4, 287.2]. At the fixed 30-FPS cadence, PPO’s three means correspond to 9.31, 8.76, and 2.13 seconds of survival. The Level-3 gain is only 18.6 frames (PPO mean 95% interval [58.8, 69.5]). Level 3 is harder and introduces yellow orbs requiring a second timed input while airborne, so these per-level results are not a controlled test of generalization. All three live-evaluation levels occur in the offline corpus, but the controller receives no live-game interaction during PPO; the results therefore measure deployment transfer from offline data and imagined policy optimization, not held-out-level generalization. Separate intervention runs suggest that skipping a few recurring blocking obstacles lets the policy reach later sections, but those assisted runs are not part of the controlled statistics.

5.2 Capture-to-action latency

End-to-end latency averages 16.289 ms (median 16.308 ms, p95 18.722 ms), and only 2 of 5,000 frames exceed the 33.3 ms budget. The reciprocal mean is about 61 FPS of compute-equivalent throughput; the deployed rate remains 30 FPS to match the training distribution.

Table 3: Optimized RTX-2060 deployment timing. Component values are means in milliseconds; end-to-end is measured independently and includes dispatch overhead.

Component	ms	Component	ms
Screen capture	2.2	Crop	0.5
Grayscale	0.5	Sobel	2.5
Downscale	2.2	FSQ encoder	2.1
Transformer	4.9	Controller	0.5
Component sum	15.3	End-to-end mean	16.3
p95	18.7	Frames over 33.3 ms	2/5,000

These timings differ from diagnostic evaluator logs that move context through CPU/NumPy and synchronize after each GPU stage. We use that evaluator for survival only and the optimized path for deployment latency.

5.3 Action-conditioned continuations

World-model generation is not capped by the 45-step PPO horizon: each predicted grid and human-selected action can be fed back into the rolling context, so the interactive loop can continue for an arbitrary number of steps unless predicted death or an optional user limit stops it. Recorded rollouts often remain visually and mechanically plausible for long stretches, including height transitions and avatar transformations, before drift produces repeated motifs, empty segments, or impossible

geometry. We therefore treat this as qualitative, open-ended level continuation rather than indefinitely coherent generation. Additional successes and failure cases are available on the [project page](#).

MIRA emphasizes that action conditioning does not itself establish action adherence and evaluates commanded-action recoverability from generated video [8]. Our matched-prefix test is a smaller qualitative intervention check rather than a population-level adherence metric. Figure 2 tests action sensitivity while holding the visual context fixed. The jump-conditioned branch clears the spike and lands; the idle-conditioned branch reaches a death score of 0.74 at $t + 5$ and terminates. Decoding is used only for inspection, not controller training.




Figure 2: Greedy continuations from one reconstructed four-frame prefix; only the first future action differs. At $t + 1, t + 5, t + 9, t + 13$, the jump branch clears the spike. The idle branch terminates at $t + 5$ (death score 0.74), so later frames are not generated.

5.4 Design diagnostics

Table 4 condenses engineering variants explored before freezing the system. Except for the quantified BC timing comparison, these are local selection diagnostics rather than balanced statistical ablations; they explain the shipped design but do not establish general architectural rankings. In particular, the joint variant is an FSQ-token adaptation inspired by LeWorldModel, not a reproduction of its continuous-latent JEPA or SIGReg regularization [17].

Table 4: Pre-freeze design diagnostics. “Live” denotes real-game progress in local evaluation.

Variant	Observed outcome	Decision
PPO without BC	Same latent plateau, about 2,000 more iterations (≈ 6 A100 hours); BC takes about 5 minutes.	Keep BC.
Joint FSQ+transformer	Lower validation CE but worse live progress; removing the pixel anchor collapsed the encoder.	Freeze tokenizer.
AdaLN action conditioning [23]	Canonical zero-gate initialization stalled; bias-1 gates plus QK-normalization [24] trained stably and reached the best development CPC. The trial co-varied other components and did not establish a live advantage over action tokens.	Use action tokens.
Transformer width 512 MLP on $\mathbf{h}_t$ only	Better CE/contrastive metrics, no live gain, slower training. PPO was 2–5 $\times$ faster with a smaller train-eval gap, but live progress was worse, suggesting lost grid detail.	Use width 384. Keep CNN.

6 Limitations

DashVMC is a single-game study with binary actions, deterministic layouts, and one training seed. The hundreds of live attempts quantify deployment variability for one frozen policy, not uncertainty across training runs, and survival time is not level completion. The BC control is an epoch-9 rerun

close to the archived training trace, not the missing original checkpoint. Because no other world-model agent is available under the same captured-pixel, non-paused Geometry-Dash protocol, the results are a controlled within-system comparison rather than a state-of-the-art benchmark.

The Geometry-Dash evidence for Equation 1 is a single-run, same-architecture comparison against uniform smoothing, not a replicated hard-CE ablation; it supports the local choice but neither quantifies uncertainty nor isolates transfer to the final system. In a matched IRIS/Pong diagnostic with five seeds per condition, annealed SLS underperformed CE in final return (12.14 ± 10.82 versus 15.68 ± 4.96) and doubled the tail-failure rate (40% versus 20%), despite lower drawdown and return volatility and a transient advantage during epochs 250–420. We therefore treat any benefit as specific to the structured FSQ geometry and Geometry-Dash setting rather than claim a general discrete-world-model improvement. The open-ended decoded continuations are qualitative: coherence is not scored, and autoregressive drift can yield repetitions, empty stretches, or impossible states; they are neither editor-valid nor guaranteed playable.

Finally, deployment is local but training is not resource-light: the modules were trained sequentially on an A100, then optimized for an RTX 2060. The 16.3-ms result covers this hardware and 64-pixel edge observations; RGB input, stochastic environments, larger action spaces, or weaker devices require new accuracy–latency measurements.

7 Conclusion

DashVMC shows that a compact discrete world model can support both imagined policy improvement and live reactive control in a timing-sensitive game. Its frozen PPO controller improves substantially over BC on the first two tested levels, while the capture-to-action path remains below the 30-FPS frame budget on a consumer GPU. Matched continuations confirm that the model’s visual futures respond to action interventions, and interactive rollouts show that those dynamics can produce open-ended, often plausible level continuations, albeit without indefinite coherence guarantees. The evidence is intentionally scoped: this is a deployed systems case study, not a general world-model benchmark or an isolated validation of structured smoothing.

Acknowledgments and Disclosure of Funding

We thank Maël Planchot for contributing gameplay data. The NVIDIA A100 GPU used for primary training was provided by MesoNet (CRIANN, Rouen) through the Representation Learning course at INSA Rouen Normandy. Geometry Dash is developed by RobTop Games.

References

- [1] David Ha and Jürgen Schmidhuber. World models. In *Neural Information Processing Systems*, 2018.
- [2] Vincent Micheli, Eloi Alonso, and François Fleuret. Transformers are sample-efficient world models. In *International Conference on Learning Representations*, 2023.
- [3] Eloi Alonso, Adam Jelley, Vincent Micheli, Anssi Kanervisto, Amos Storkey, Tim Pearce, and François Fleuret. Diffusion for world modeling: Visual details matter in atari. In *Advances in Neural Information Processing Systems*, volume 37, 2024.
- [4] Fabian Mentzer, David Minnen, Eiríkur Agustsson, and Michael Tschannen. Finite scalar quantization: VQ-VAE made simple. In *International Conference on Learning Representations*, 2024.
- [5] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.
- [6] Vincent Micheli, Eloi Alonso, and François Fleuret. Efficient world models with context-aware tokenization. In *International Conference on Machine Learning*, 2024.
- [7] Danijar Hafner, Jurgis Pasukonis, Jimmy Ba, and Timothy Lillicrap. Mastering diverse control tasks through world models. *Nature*, 2025.

- [8] Anthony Hu, Václav Volhejn, Adrien Ramanana Rahary, Chris Mulder, Aditya Makkar, Alyx Liao, Amélie Royer, Manu Orsini, Adam Jelley, Eloi Alonso, Florian Laurent, Fredrik Norén, James Swingos, Jan Hünemann, Kent Rollins, Lucas Hosseini, Matthieu Le Cauchois, Maxim Peter, Pim de Witte, Tim Brown, Vincent Micheli, Moritz Böhle, Gabriel de Marmiesse, Viktoriia Sharmanska, Lucia Specia, Michael Black, and Patrick Pérez. Multiplayer interactive world models with representation autoencoders. *arXiv preprint arXiv:2607.05352*, 2026.
- [9] Maxime Burchi and Radu Timofte. TWISTER: Learning transformer-based world models with contrastive predictive coding. In *International Conference on Learning Representations*, 2025.
- [10] Junliang Guo, Yang Ye, Tianyu He, Haoyu Wu, Yushu Jiang, Tim Pearce, and Jiang Bian. MineWorld: A real-time and open-source interactive world model on Minecraft. *arXiv preprint arXiv:2504.08388*, 2025.
- [11] Xianglong He, Chunli Peng, Zexiang Liu, Boyang Wang, Yifan Zhang, Qi Cui, Fei Kang, Biao Jiang, Mengyin An, Yangyang Ren, Baixin Xu, Hao-Xiang Guo, Kaixiong Gong, Size Wu, Wei Li, Xuchen Song, Yang Liu, Yangguang Li, and Yahui Zhou. Matrix-Game 2.0: An open-source real-time and streaming interactive world model. *arXiv preprint arXiv:2508.13009*, 2025.
- [12] Ted Li and Sean Rafferty. Playing Geometry Dash with convolutional neural networks. Stanford CS231N course report, 2017.
- [13] Bin Lin, Zongjian Li, Yuwei Niu, Kaixiong Gong, Yunyang Ge, Yunlong Lin, Mingzhe Zheng, Jianwei Zhang, Miles Yang, Zhao Zhong, Liefeng Bo, and Li Yuan. iFSQ: Improving FSQ for image generation with 1 line of code. *arXiv preprint arXiv:2601.17124*, 2026.
- [14] Todd Hester, Matej Vecerik, Olivier Pietquin, Marc Lanctot, Tom Schaul, Bilal Piot, Dan Horgan, John Quan, Andrew Sendonaris, Gabriel Dulac-Arnold, Ian Osband, John Agapiou, Joel Z. Leibo, and Audrunas Gruslys. Deep q-learning from demonstrations. In *AAAI Conference on Artificial Intelligence*, 2018.
- [15] Bowen Baker, Ilge Akkaya, Peter Zhokhov, Joost Huizinga, Jie Tang, Adrien Ecoffet, Brandon Houghton, Raul Sampedro, and Jeff Clune. Video pretraining (VPT): Learning to act by watching unlabeled online videos. *arXiv preprint arXiv:2206.11795*, 2022.
- [16] Hugues Van Assel, Mark Ibrahim, Tommaso Biancalani, Aviv Regev, and Randall Balestriero. Joint embedding vs reconstruction: Provable benefits of latent space prediction for self supervised learning. *arXiv preprint arXiv:2505.12477*, 2025.
- [17] Lucas Maes, Quentin Le Lidec, Damien Scieur, Yann LeCun, and Randall Balestriero. LeWorld-Model: Stable end-to-end joint-embedding predictive architecture from pixels. *arXiv preprint arXiv:2603.19312*, 2026.
- [18] Tongda Xu, Wendi Zheng, Jiajun He, José Miguel Hernández-Lobato, Yan Wang, Ya-Qin Zhang, and Jie Tang. Training-free vector quantization via gaussian VAEs. *arXiv preprint arXiv:2512.06609*, 2025.
- [19] Zaishuo Xia, Yukuan Lu, Xinyi Li, Yifan Xu, and Yubei Chen. Clone deterministic 3d worlds with geometrically-regularized world models. *arXiv preprint arXiv:2510.26782*, 2025.
- [20] Jianlin Su, Murtadha Ahmed, Yu Lu, Shengfeng Pan, Wen Bo, and Yunfeng Liu. Roformer: Enhanced transformer with rotary position embedding. *Neurocomputing*, 568:127063, 2024.
- [21] Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. Focal loss for dense object detection. In *International Conference on Computer Vision*, pages 2980–2988, 2017.
- [22] Christian Szegedy, Vincent Vanhoucke, Sergey Ioffe, Jon Shlens, and Zbigniew Wojna. Rethinking the Inception architecture. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 2818–2826, 2016.
- [23] William Peebles and Saining Xie. Scalable diffusion models with transformers. In *International Conference on Computer Vision*, 2023.

- [24] Patrick Esser, Sumith Kulal, Andreas Blattmann, Rahim Entezari, Jonas Müller, Harry Saini, Yam Levi, Dominik Lorber, Axel Sauer, Frederic Boesel, et al. Scaling rectified flow transformers for high-resolution image synthesis. In *International Conference on Machine Learning*, 2024.